

HW 3: Ensemble Techniques

Atharva Bhagwat

December 5, 2021

CS-GY 6923: MACHINE LEARNING
Professor: Dr. Raman Kannan

Contents

1	Overview of the Dataset	3
2	Loading the Dataset	4
3	Exploratory Data Analysis	6
3.1	Understanding the Data	6
3.1.1	Target Variable	6
3.1.2	Types of Features	7
3.2	Feature Analysis	7
3.2.1	Categorical Features	7
3.2.2	Continuous Features	11
3.2.3	Missing Values	14
3.2.4	Handling Missing Values	15
3.3	Correlation	15
3.4	Eliminating Correlated Features	16
3.5	Variable Importance using Random Forest Classifier	17
3.6	Variable Inflation Factor	17
4	EDA: Conclusion	19
4.1	Observations	19
5	Implementing Classifiers	20
5.1	GLM	20
5.1.1	P-Value Analysis	20
5.1.2	Confusion Matrix	20
5.2	Naive Bayes	23
5.2.1	Confusion Matrix	23
5.3	Support Vector Machines	25
5.3.1	Confusion Matrix	25
5.4	Decision Tree	27
5.4.1	Confusion Matrix	27
5.5	K Nearest Neighbor	30
5.5.1	Confusion Matrix	30
5.6	Tabulating Model Performance on Test Set	32

5.7	Observations:	32
6	Ensemble Techniques	34
6.1	K Fold Cross Validation	34
6.2	Boosting	35
6.3	Random Forest	36
6.4	Tabulating Ensemble Method Performance	39
6.5	Conclusion	39

Chapter 1

Overview of the Dataset

The dataset is called "[Airlines Codrna Adult](#)" and was taken from OpenML.org.

The dataset consists of 29 independent variables and 1 dependent variable and has 1,076,790 observations. The dependent variable is called 'Delay' and is a categorical variable with only two values, 0 or 1. Therefore, this is a binary classification problem. The dataset is a CSV file and will be worked on using R.

Chapter 2

Loading the Dataset

The dataset CSV file is named 'AirlinesCodrnaAdult.csv' and is located inside 'dataset' folder. It is loaded into a variable named *data*.

```
# Function to Load Dataset
load_dataset = function(csvfile){read.csv(csvfile)}

cat("Loading Dataset...")
data = load_dataset('/home/2021/nyu/fall/acb9244/dataset/AirlinesCodrnaAdult.csv")      # Load Data
```

We will perform the following checks to see if the data is loaded properly.

1. Display the first few entries using the command `head(data)`.
2. Check number of observations and number of features using `dim(data)`.
3. Display the column names using the command `names(data)`.

```
[> head(data)
  Airline Flight AirportFrom AirportTo DayOfWeek Time Length codrna_X1
1   US    1558      PHX       CLT      3 15 222 0.698608
2   AA    2400      LAX       DFW      3 20 165 0.698608
3   AA    2466      SFO       DFW      3 20 195 0.698608
4   AS     108      ANC       SEA      3 30 202 0.698608
5   CO    1094      LAX       IAH      3 30 181 0.698608
6   DL    1768      LAX       MSP      3 30 220 0.698608
codrna_X2 codrna_X3 codrna_X4 codrna_X5 codrna_X6 codrna_X7 codrna_X8 age
1 0.136363 0.281746 0.286873 0.599079 0.204906 0.569492 0.587408 2
2 0.136363 0.281746 0.286873 0.599079 0.204906 0.569492 0.587408 2
3 0.136363 0.281746 0.286873 0.599079 0.204906 0.569492 0.587408 2
4 0.136363 0.281746 0.286873 0.599079 0.204906 0.569492 0.587408 2
5 0.136363 0.281746 0.286873 0.599079 0.204906 0.569492 0.587408 2
6 0.136363 0.281746 0.286873 0.599079 0.204906 0.569492 0.587408 2
  workclass fnlwgt education education.num marital.status occupation
1 State-gov 77516 Bachelors          13 Never-married Adm-clerical
2 State-gov 77516 Bachelors          13 Never-married Adm-clerical
3 State-gov 77516 Bachelors          13 Never-married Adm-clerical
4 State-gov 77516 Bachelors          13 Never-married Adm-clerical
5 State-gov 77516 Bachelors          13 Never-married Adm-clerical
6 State-gov 77516 Bachelors          13 Never-married Adm-clerical
  relationship race sex capitalgain capitalloss hoursperweek native.country
1 Not-in-family White Male           1          0          2 United-States
2 Not-in-family White Male           1          0          2 United-States
3 Not-in-family White Male           1          0          2 United-States
4 Not-in-family White Male           1          0          2 United-States
5 Not-in-family White Male           1          0          2 United-States
6 Not-in-family White Male           1          0          2 United-States
  Delay
1 1
2 1
3 1
4 0
5 1
6 0
[> nrow(data)
[1] 1076790
[> ncol(data)
[1] 30
[> names(data)
 [1] "Airline"         "Flight"        "AirportFrom"     "AirportTo"
 [5] "DayOfWeek"       "Time"          "Length"         "codrna_X1"
 [9] "codrna_X2"       "codrna_X3"     "codrna_X4"       "codrna_X5"
 [13] "codrna_X6"       "codrna_X7"     "codrna_X8"       "age"
 [17] "workclass"       "fnlwgt"        "education"      "education.num"
 [21] "marital.status"  "occupation"    "relationship"   "race"
 [25] "sex"             "capitalgain"  "capitalloss"    "hoursperweek"
 [29] "native.country"  "Delay"
```

We see that the dataset contains 1,076,790 observations and 30 features. We will further analyse the data in the next chapter.

Chapter 3

Exploratory Data Analysis

3.1 Understanding the Data

The first step to building machine learning models is understanding the data.

3.1.1 Target Variable

We will start with the target variable. We will look at number of classes, number of observations for each class and their percentage.

```
target_information <- function(dataframe){  
  dataframe.target_unique_val <- unique(data$Delay)  
  cat('Class values in "Delay" column, ie: target variable',dataframe.target_unique_val)  
  
  dataframe.target_unique_count <- table(data$Delay)  
  cat('Number of observations in each class',dataframe.target_unique_count)  
  
  dataframe.target_unique_per <- (dataframe.target_unique_count/nrow(data))*100  
  cat('Percent of observations in each class',dataframe.target_unique_per)  
}  
  
[> target_information()  
  
Class values in "Delay" column, ie: target variable  
[1] 1 0  
Levels: 0 1  
  
Number of observations in each class  
  
      0      1  
473652 603138  
  
Percent of observations in each class  
  
      0      1  
43.98741 56.01259
```

The function `target_information()`, prints the unique class values, number of

observations for each class and their percentage. As we see only two values 0 and 1, we can say that this is a binary classification problem. We also observe that the classes are slightly imbalanced.

3.1.2 Types of Features

We have total of 29 features. A feature can be categorical or continuous. We can distinguish between such features just by looking at the values they hold.

Categorical Features: Airline, Flight, AirportFrom, AirportTo, DayOfWeek, age, workclass, education, education.num, marital.status, occupation, relationship, race, sex, hoursperweek, native.country.

Continuous Features: Time, Length, codrna_X1, codrna_X2, codrna_X3, codrna_X4, codrna_X5, codrna_X6, codrna_X7, codrna_X8, fnlwgt, capitalgain, capitalloss.

We can look at the datatype of all the features using `sapply(data,class)`.

Columns and their class:					
Airline	Flight	AirportFrom	AirportTo	DayOfWeek	
"factor"	"integer"	"factor"	"factor"	"integer"	
Time	Length	codrna_X1	codrna_X2	codrna_X3	
"integer"	"integer"	"numeric"	"numeric"	"numeric"	
codrna_X4	codrna_X5	codrna_X6	codrna_X7	codrna_X8	
"numeric"	"numeric"	"numeric"	"numeric"	"numeric"	
age	workclass	fnlwgt	education	education.num	
"integer"	"factor"	"integer"	"factor"	"integer"	
marital.status	occupation	relationship	race	sex	
"factor"	"factor"	"factor"	"factor"	"factor"	
capitalgain	capitalloss	hoursperweek	native.country	Delay	
"integer"	"integer"	"integer"	"factor"	"factor"	

3.2 Feature Analysis

We learnt that our features are of two types: categorical and continuous. A categorical feature is variable that contain a finite number of unique values. Whereas a continuous features have infinite values between two values.

3.2.1 Categorical Features

Tests to run:

1. Check number of unique values.
2. Plot bar graphs to look at the distribution/percentage of unique values on target variable

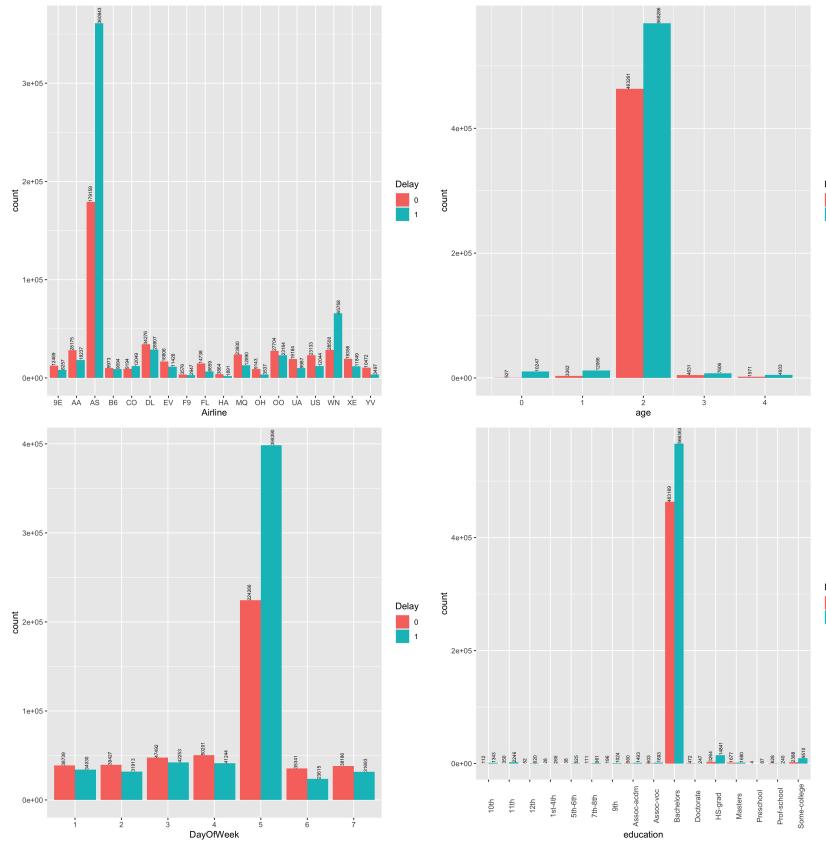
```

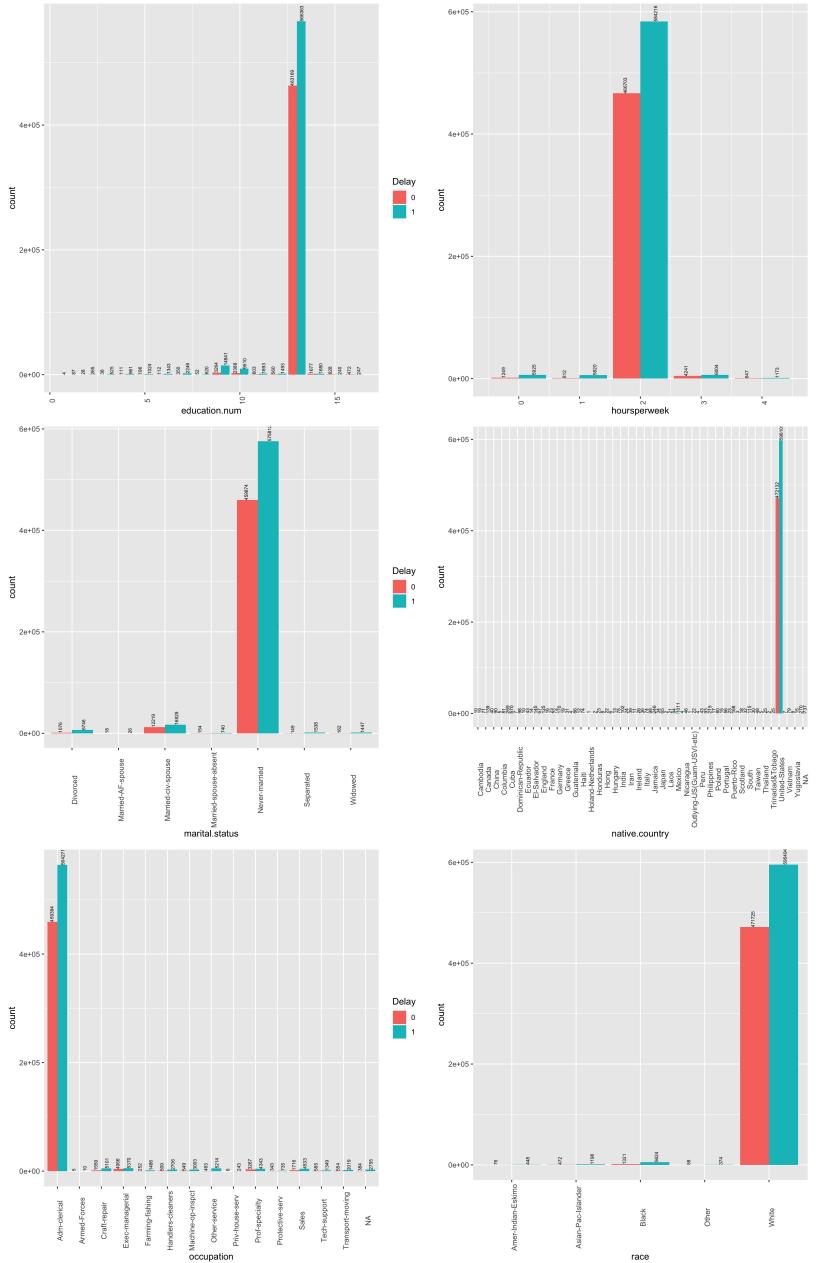
> unique_values <- function(){
+   for (column in categorical_features){
+     pprint(paste('Unique values in',column), table(data[column]))
+   }
+ }
> probability_dist <- function(){
+   for (column in categorical_features){
+     pprint(paste('Percentage of unique values in',column),round(table(data[c(column,'Delay')])*100/nrow(data),digits=2))
+   }
+ }

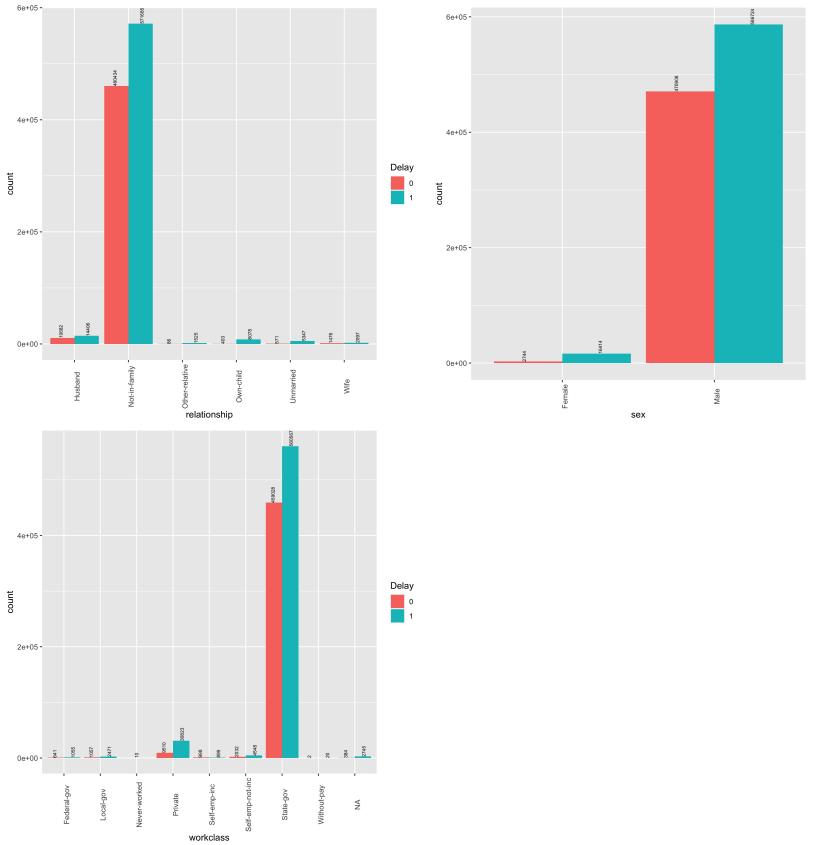
```

We see that for every feature a single category has very large number of observations.

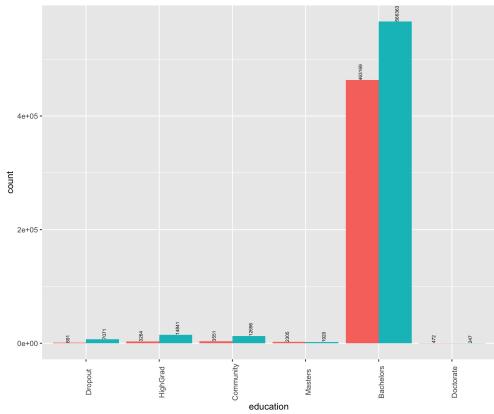
Distribution for character features:







From the above graphs we see that 'education' and 'education.num' deliver the same information. So we will drop column 'education.num' using `data[, !names(data) %in% c('education.num')]`. We will also see that education has 16 unique values. We will recast this feature into broader groups using `recode_factor()`. We do this to group similar categories together to extract more information.



Old Label	New Label
Preschool	Dropout
1 st -4 th	Dropout
5 th -6 th	Dropout
7 th -8 th	Dropout
9 th	Dropout
10 th	Dropout
11 th	Dropout
12 th	Dropout
HS-grad	HighGrad
Some-college	Community
Assoc-acdm	Community
Assoc-voc	Community
Bachelors	Bachelors
Masters	Masters
Prof-school	Masters
Doctorate	Doctorate

3.2.2 Continuous Features

Tests to run:

1. Calculate mean, median, mode, minimum, maximum, standard deviation, 1st and 3rd quartiles.
2. Check for outliers.

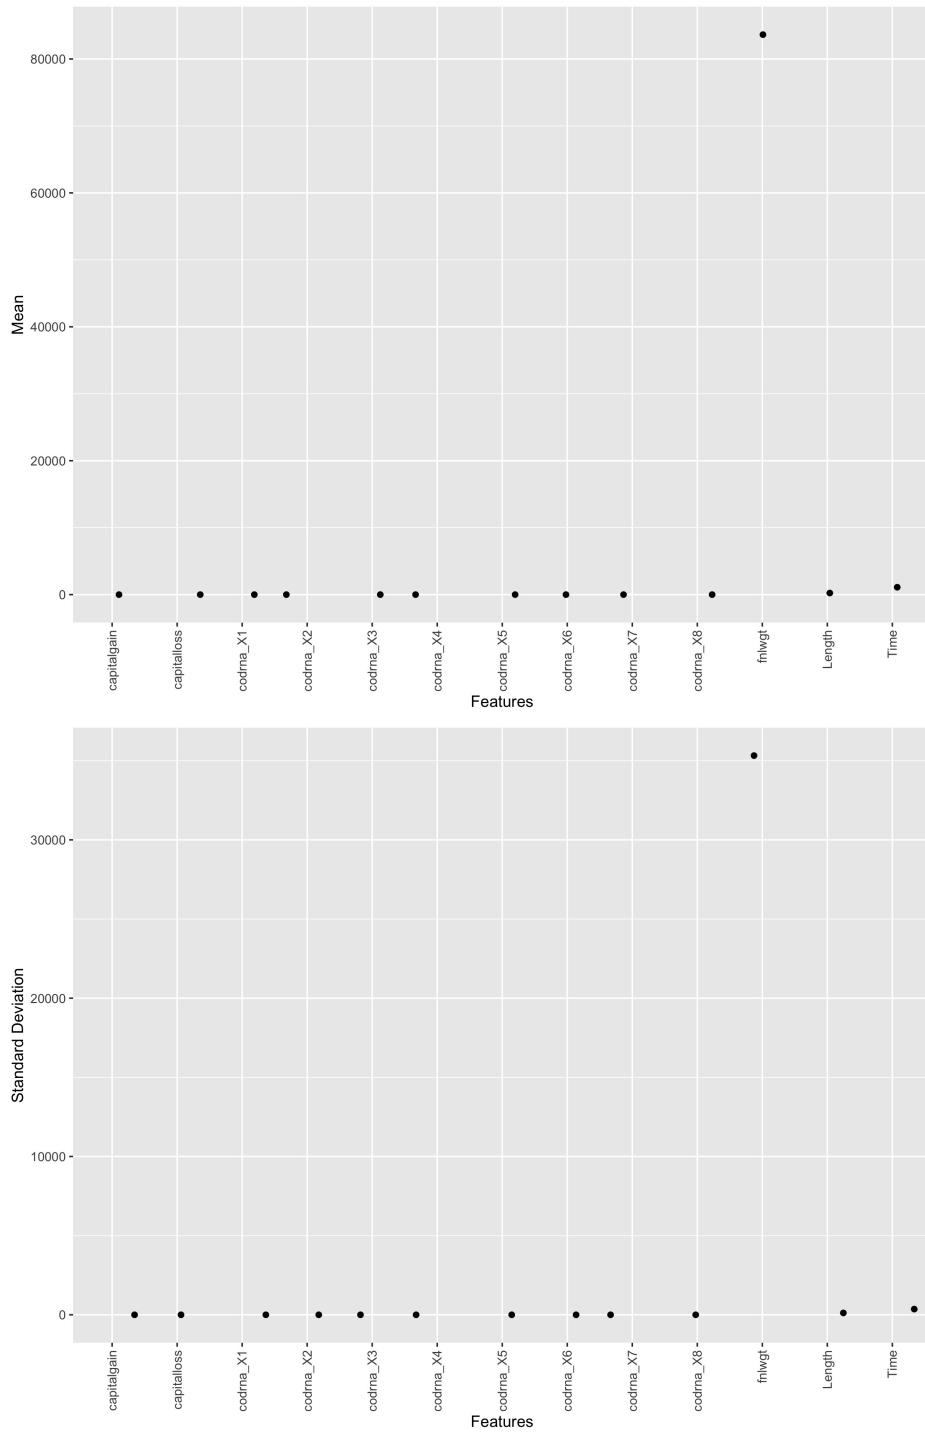
Many of these properties can be calculated using `summary(data[continuous_features])`. We calculate the outliers using the range $[Q1 - IQR * 1.5, Q3 + IQR * 1.5]$.

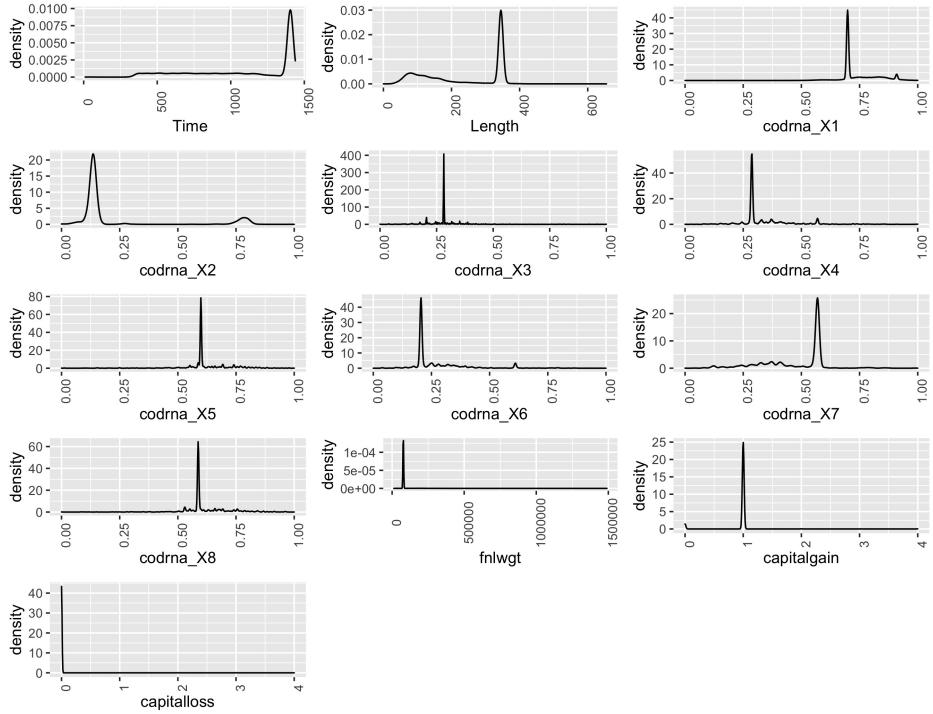
```

Summary:
      Time      Length     codrna_X1     codrna_X2
Min. : 10    Min. : 0.0    Min. :0.0000    Min. :0.0000
1st Qu.: 795  1st Qu.:115.0   1st Qu.:0.6986   1st Qu.:0.1364
Median :1355  Median :345.0   Median :0.6986   Median :0.1364
Mean   :1103  Mean   :237.2   Mean   :0.7453   Mean   :0.2219
3rd Qu.:1405 3rd Qu.:345.0   3rd Qu.:0.8009   3rd Qu.:0.1364
Max.  :1439   Max.  :655.0   Max.  :1.0000   Max.  :1.0000
      codrna_X3     codrna_X4     codrna_X5     codrna_X6
Min. :0.0000   Min. :0.0000   Min. :0.0000   Min. :0.0000
1st Qu.:0.2744 1st Qu.:0.2869  1st Qu.:0.5991  1st Qu.:0.2049
Median :0.2817  Median :0.2869  Median :0.5991  Median :0.2049
Mean   :0.2907  Mean   :0.3410  Mean   :0.6204  Mean   :0.2709
3rd Qu.:0.2817 3rd Qu.:0.3711  3rd Qu.:0.6590  3rd Qu.:0.3080
Max.  :1.0000   Max.  :1.0000  Max.  :1.0000  Max.  :1.0000
      codrna_X7     codrna_X8     fnlwgt      capitalgain
Min. :0.0000   Min. :0.0000   Min. : 12285  Min. :0.0000
1st Qu.:0.3711 1st Qu.:0.5874  1st Qu.: 77516  1st Qu.:1.0000
Median :0.5695  Median :0.5874  Median : 77516  Median :1.0000
Mean   :0.4763  Mean   :0.6145  Mean   : 83638  Mean   :0.9551
3rd Qu.:0.5695 3rd Qu.:0.6590  3rd Qu.: 77516  3rd Qu.:1.0000
Max.  :1.0000   Max.  :1.0000  Max.  :1490400  Max.  :4.0000
      capitalloss
Min. :0.000000
1st Qu.:0.000000
Median :0.000000
Mean   :0.005733
3rd Qu.:0.000000
Max.  :4.000000

Standard deviation:
      Time      Length     codrna_X1     codrna_X2     codrna_X3     codrna_X4
      359.38    117.41     0.08      0.22      0.11      0.12
      codrna_X5     codrna_X6     codrna_X7     codrna_X8     fnlwgt capitalgain
      0.11      0.13      0.14      0.11    35324.01      0.25
      capitalloss
      0.13

```





```

> data.outliers <- c()
> for(column in continuous_features){
+   data_mean <- mean(data[,column])
+   data_sd <- sd(data[,column])
+   lower_cutoff <- data_mean-3*data_sd
+   higher_cutoff <- data_mean+3*data_sd
+   data.outliers_id <- which(data[,column]<lower_cutoff | data[,column]>higher_cutoff)
+   data.outliers <- c(data.outliers,data.outliers_id)
+ }
> data.outliers <- unique(data.outliers)
> pprint('Number of Outliers:',length(data.outliers))

Number of Outliers:
[1] 133877
> pprint('Percentage of Outliers:',round(length(data.outliers)*100/nrow(data),2))

Percentage of Outliers:
[1] 12.43

```

All continuous features have single peak. 12.43% of the total observations are outliers.

3.2.3 Missing Values

We inspect the data to see if there are any missing values.

```

missing_value_info <- function(dataframe){
  dataframe.missing_count <- sapply(dataframe, function(x) sum(is.na(x)))
  pprint('Number of missing values:',dataframe.missing_count)

  dataframe.missing_per <- (dataframe.missing_count/nrow(dataframe))*100
  pprint('Percentage of missing values:',dataframe.missing_per)
}

[> missing_value_info(data)

Number of missing values:
  Airline      Flight    AirportFrom    AirportTo   DayOfWeek
  0            0           0             0           0
  Time        Length    codrna_X1    codrna_X2   codrna_X3
  0            0           0             0           0
codrna_X4    codrna_X5    codrna_X6    codrna_X7   codrna_X8
  0            0           0             0           0
  age        workclass   fnlwgt    education  education.num
  0            3129       0             0           0
marital.status occupation relationship race       sex
  0            3139       0             0           0
  capitalgain capitalloss hoursperweek native.country Delay
  0            0           0             0           1007          0

Percentage of missing values:
  Airline      Flight    AirportFrom    AirportTo   DayOfWeek
  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
  Time        Length    codrna_X1    codrna_X2   codrna_X3
  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
codrna_X4    codrna_X5    codrna_X6    codrna_X7   codrna_X8
  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
  age        workclass   fnlwgt    education  education.num
  0.0000000  0.2905859  0.0000000  0.0000000  0.0000000
marital.status occupation relationship race       sex
  0.0000000  0.2915146  0.0000000  0.0000000  0.0000000
  capitalgain capitalloss hoursperweek native.country Delay
  0.0000000  0.0000000  0.0000000  0.0935187  0.0000000

```

We see that 3 features: workclass (0.29%), occupation (0.29%), and native.country(0.09%) have missing values.

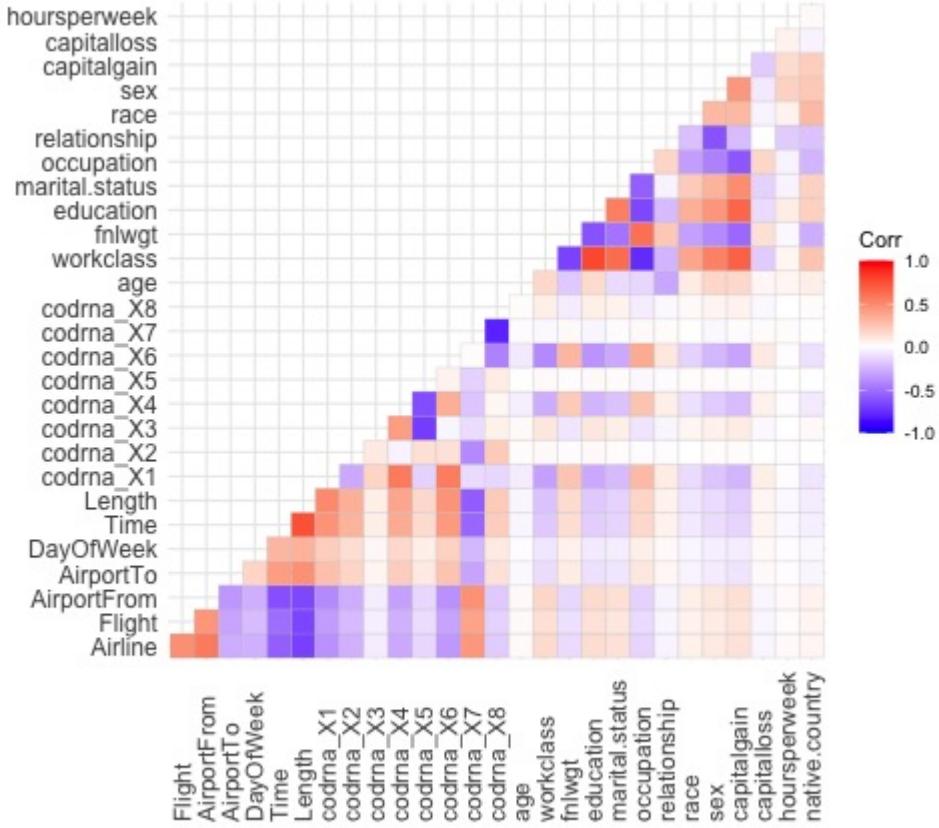
3.2.4 Handling Missing Values

Only three features: work class, occupation, and native country have missing values. The missing values are so small in number that we can just drop the observation with missing values using `drop_na(data)`.

3.3 Correlation

As observed in [figure](#) not all features are of type numeric. We first convert these features to numeric so we can perform further steps. We use `as.numeric(data[column])` to convert the features to numeric.

We now calculate the correlation using `cor(data)`.



We plot the lower correlation matrix. Correlation values equal to 0 means that the variables are not correlated and values 1 and -1 mean that the variable are positively and negatively correlated respectively. After checking how many of the feature are highly correlated, ie: the absolute correlation coefficient > 0.5 , we find 14 correlated features.

3.4 Eliminating Correlated Features

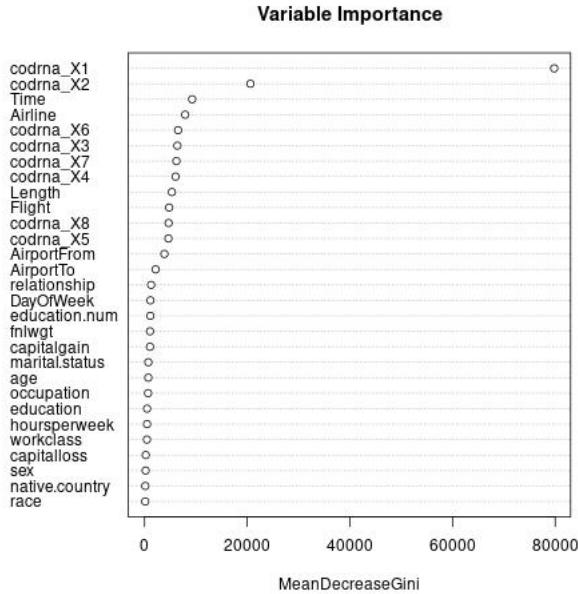
In 3.3 we saw that 14 features are highly correlated. We will eliminate these features as it would reduce the number of variables, reducing the model's complexity. We use a threshold of 0.5.

3.5 Variable Importance using Random Forest Classifier

The next step is to rank the features as per their importance. This is done using a Random Forest Classifier and the `varImp()` from the caret package.

To build a classifier, we divide the data into train and test sets. We will use 70% of data as the training set and 30% as testing test. We will randomize the dataset before splitting and will use `set.seed(42)` to replicate the result.

```
> set.seed(42)
> data.randomized <- data[sample(1:nrow(data)), nrow(data)], ]
>
> split_indexes <- sample(1:nrow(data), 0.7*nrow(data), replace=F)
> data.train <- data.randomized[split_indexes, ]
> data.test <- data.randomized[-split_indexes, ]
>
[> data.train_rf <- data.train
```



3.6 Variable Inflation Factor

VIF is a measure of multicollinearity among the variables. Multicollinearity is when two or more features are highly correlated to each other. This measure leads to higher variance in the model. To check VIF we will first convert the Delay variable to numeric as VIF can only be calculated for regression problems. We will use a generalized linear model to train.

```
Training GLM...> glm_model <- glm(Delay~., data.train_vif, family='binomial')
> glm_model_vif <- car::vif(glm_model)
[> pprint('Variable Inflation Factor in ascending order:', sort(glm_model_vif))

Variable Inflation Factor in ascending order:
  age DayOfWeek education   Flight workclass   Airline codrna_X1
1.021245 1.078048 1.171476 1.290806 1.302775 1.361159 1.414832
```

After removing the correlating features, we observe that the VIF values are not very high. This means that the chances of these set of features having high multicollinearity are low.

Chapter 4

EDA: Conclusion

4.1 Observations

- This is a binary classification problem.
- Dataset consists of both categorical and continuous features.
- Both categorical and continuous features have a single value with majority of observations.
- Feature 'education' is grouped into less number of categories to extract more information.
- Dataset contains very low number of missing values(0.6%). Hence the rows with missing values can be dropped.
- 14 features are highly correlated and should be dropped.
- Individual classifiers should be trained after correlated features are dropped and evaluated based on p-value, accuracy, recall, specificity and auc.

Chapter 5

Implementing Classifiers

The target variable for this dataset is 'Delay'. During application, it is crucial for the model to predict a +ve observation to be a +ve class. This is very similar to detection of cancer. So, for the purpose of analysis we will give preference to recall and precision values.

5.1 GLM

5.1.1 P-Value Analysis

P-values for all the features are less than 0.05. This implies that we can reject the null hypothesis. All the features have statistical relation with the target variable.

5.1.2 Confusion Matrix

The accuracy, recall(sensitivity), specificity and precision for both training and testing phase is very similar. However we do observe that the recall value for is lower than the precision. This means that the number of samples correctly classified as +ve might be low but the probability of the classification being correct is high.

```

GLM Summary:

Call:
glm(formula = Delay ~ ., family = "binomial", data = data.train)

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-3.9211 -0.9772  0.3186  0.8838  4.2328 

Coefficients:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -1.763e+01 1.392e-01 -126.65 < 2e-16 ***
Airline      7.775e-02 6.009e-04 129.38 < 2e-16 ***
Flight       3.702e-05 1.568e-06 23.62 < 2e-16 ***  
AirportTo    -1.105e-03 4.158e-05 -26.58 < 2e-16 ***  
DayOfWeek   -9.450e-02 1.765e-03 -53.55 < 2e-16 ***  
codrna_X1    2.244e+01 6.398e-02 350.69 < 2e-16 ***  
codrna_X2    4.149e+00 1.820e-02 227.99 < 2e-16 ***  
codrna_X3   -3.729e+00 3.051e-02 -122.25 < 2e-16 ***  
age          -5.638e-01 1.356e-02 -41.59 < 2e-16 ***  
workclass    4.952e-01 5.885e-03  84.14 < 2e-16 ***  
relationship 3.500e-01 1.104e-02  31.70 < 2e-16 ***  
race         -2.331e-01 2.036e-02 -11.45 < 2e-16 ***  
capitalloss  -6.575e-01 2.388e-02 -27.54 < 2e-16 ***  
hoursperweek -7.319e-01 1.745e-02 -41.95 < 2e-16 ***  
native.country 1.797e-02 2.197e-03   8.18 2.85e-16 ***  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1030307  on 750892  degrees of freedom
Residual deviance:  808650  on 750878  degrees of freedom
AIC: 808680

Number of Fisher Scoring iterations: 5

```

```

GLM Training Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 246529 138993
      1  84254 281117

      Accuracy : 0.7027
      95% CI : (0.7017, 0.7037)
      No Information Rate : 0.5595
      P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.4073

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.6692
      Specificity : 0.7453
      Pos Pred Value : 0.7694
      Neg Pred Value : 0.6395
      Prevalence : 0.5595
      Detection Rate : 0.3744
      Detection Prevalence : 0.4866
      Balanced Accuracy : 0.7072

      'Positive' Class : 1

GLM Testing Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 105842  59292
      1  36385 120293

      Accuracy : 0.7027
      95% CI : (0.7011, 0.7043)
      No Information Rate : 0.558
      P-Value [Acc > NIR] : < 2.2e-16

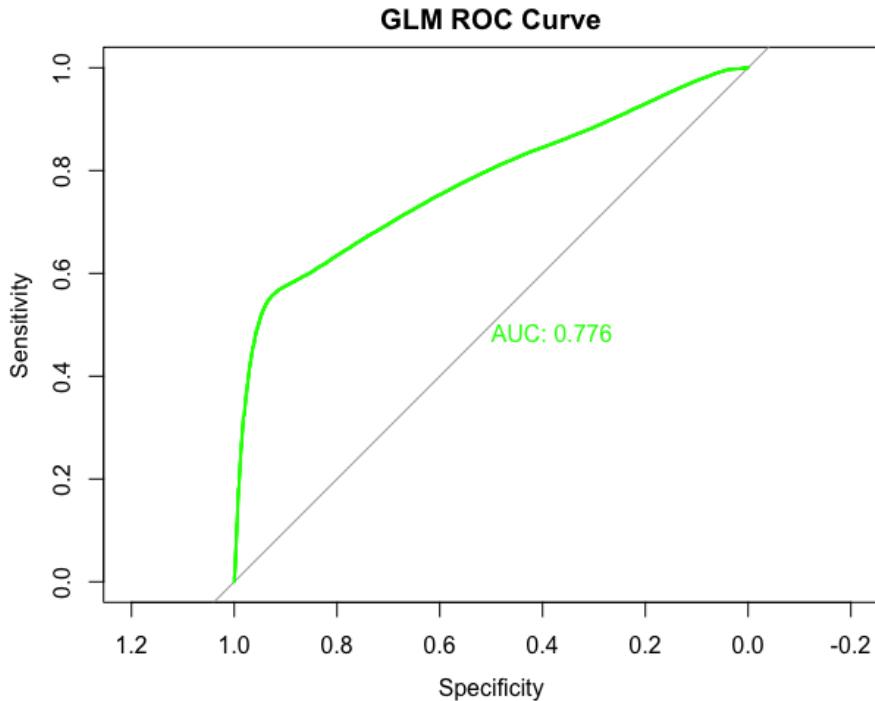
      Kappa : 0.4072

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.6698
      Specificity : 0.7442
      Pos Pred Value : 0.7678
      Neg Pred Value : 0.6409
      Prevalence : 0.5580
      Detection Rate : 0.3738
      Detection Prevalence : 0.4869
      Balanced Accuracy : 0.7070

      'Positive' Class : 1

```



5.2 Naive Bayes

5.2.1 Confusion Matrix

For both the training and testing phase we see accuracy score, precision, recall and specificity scores to be very similar. But, recall value is too low and precision and specificity is high. Meaning the model is returning a positive class for only sample it is confident about and returning negative for the rest. Model has low accuracy of detecting +ve samples.

```

Naive Bayes Training Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 304848 263862
      1 25935 156248

      Accuracy : 0.6141
      95% CI : (0.613, 0.6152)
      No Information Rate : 0.5595
      P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.2727

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.3719
      Specificity : 0.9216
      Pos Pred Value : 0.8576
      Neg Pred Value : 0.5360
      Prevalence : 0.5595
      Detection Rate : 0.2081
      Detection Prevalence : 0.2426
      Balanced Accuracy : 0.6468

      'Positive' Class : 1

Naive Bayes Testing Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 131047 112737
      1 11180 66848

      Accuracy : 0.6149
      95% CI : (0.6133, 0.6166)
      No Information Rate : 0.558
      P-Value [Acc > NIR] : < 2.2e-16

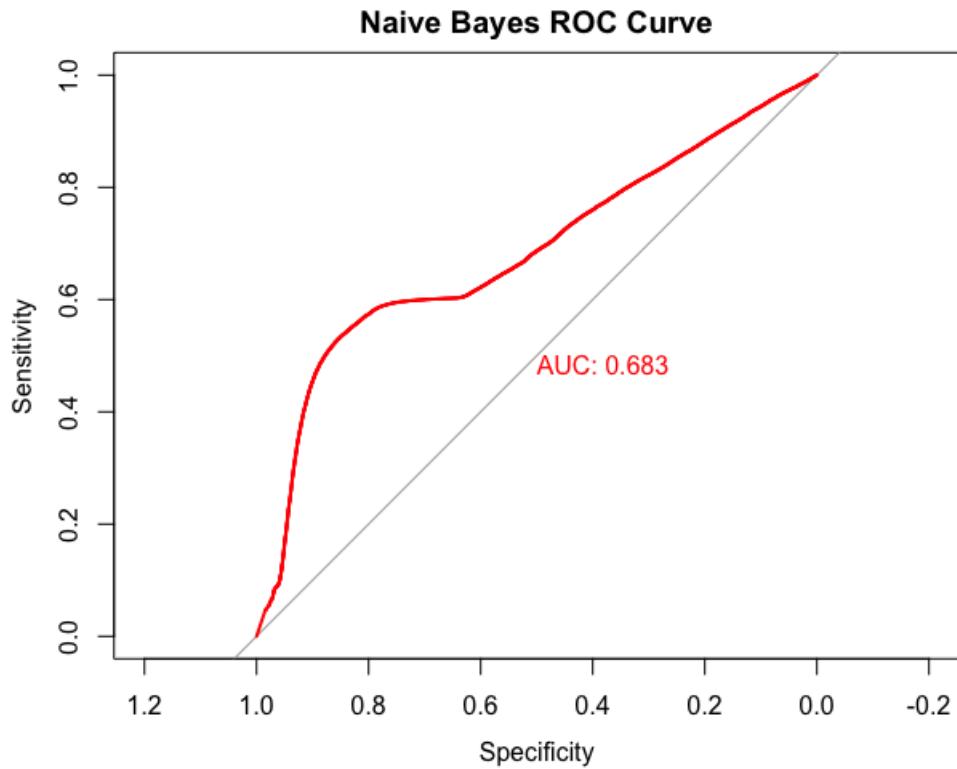
      Kappa : 0.2733

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.3722
      Specificity : 0.9214
      Pos Pred Value : 0.8567
      Neg Pred Value : 0.5376
      Prevalence : 0.5580
      Detection Rate : 0.2077
      Detection Prevalence : 0.2425
      Balanced Accuracy : 0.6468

      'Positive' Class : 1

```



5.3 Support Vector Machines

5.3.1 Confusion Matrix

For both the training and testing phase we see accuracy score, precision, recall and specificity scores to be very similar. However, similar to GLM we observe that the recall value for is lower than the precision. SVM has a better precision value than GLM. It implies that the +ve observations correctly classified by SVM have a higher chance of actually being a +ve observation.

```

SVM Testing Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 121495  61101
      1 20732   118484

      Accuracy : 0.7457
      95% CI  : (0.7442, 0.7472)
      No Information Rate : 0.558
      P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.4993

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.6598
      Specificity : 0.8542
      Pos Pred Value : 0.8511
      Neg Pred Value : 0.6654
      Prevalence : 0.5580
      Detection Rate : 0.3682
      Detection Prevalence : 0.4326
      Balanced Accuracy : 0.7570

      'Positive' Class : 1

SVM Training Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 283007 142838
      1 47776  277272

      Accuracy : 0.7462
      95% CI  : (0.7452, 0.7471)
      No Information Rate : 0.5595
      P-Value [Acc > NIR] : < 2.2e-16

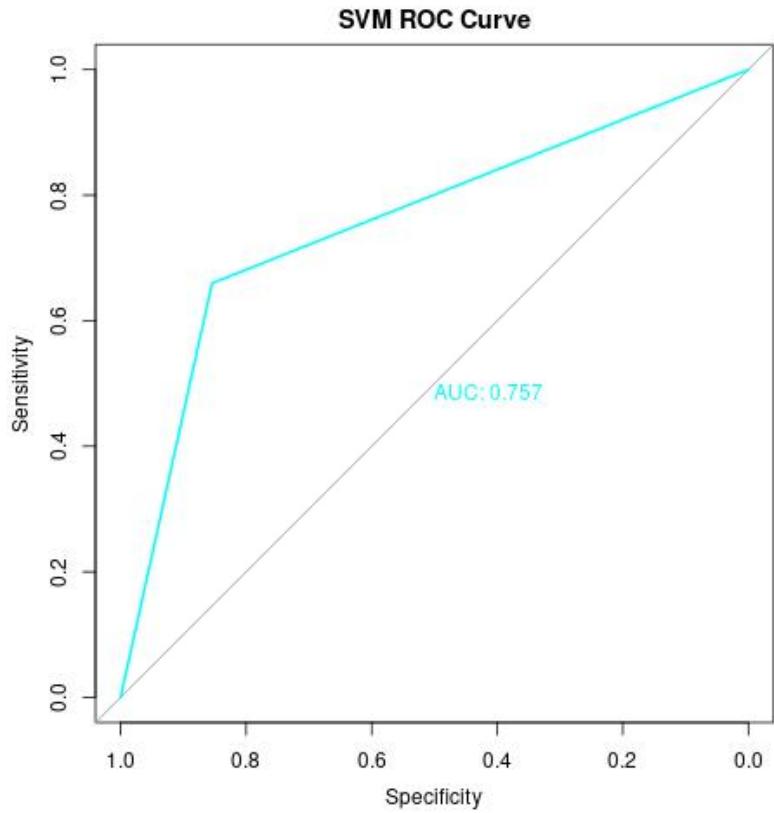
      Kappa : 0.5003

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.6600
      Specificity : 0.8556
      Pos Pred Value : 0.8530
      Neg Pred Value : 0.6646
      Prevalence : 0.5595
      Detection Rate : 0.3693
      Detection Prevalence : 0.4329
      Balanced Accuracy : 0.7578

      'Positive' Class : 1

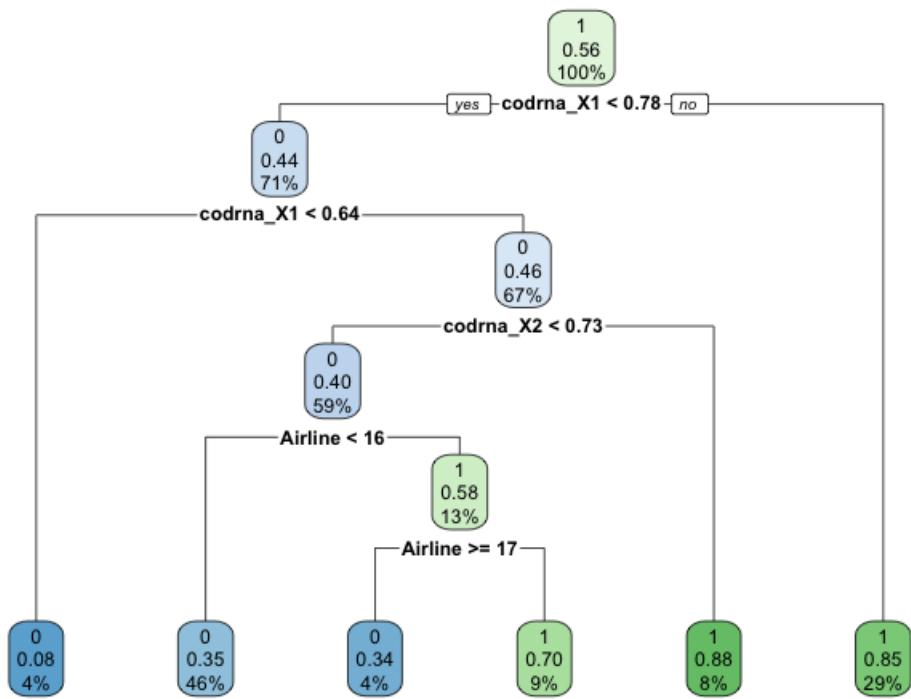
```



5.4 Decision Tree

5.4.1 Confusion Matrix

For both the training and testing phase we see accuracy score, precision, recall and specificity scores to be very similar. But we see a slight increase in testing accuracy. This is a bad sign as the model is optimized for the training set. One of the reasons that this could have happened is because of the distribution of features in training and testing set. Applying a different seed might solve this problem.



```

Decision Tree Training Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0  271424 136173
      1   59359 283937

      Accuracy : 0.7396
      95% CI : (0.7386, 0.7406)
      No Information Rate : 0.5595
      P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.4845

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.6759
      Specificity : 0.8206
      Pos Pred Value : 0.8271
      Neg Pred Value : 0.6659
      Prevalence : 0.5595
      Detection Rate : 0.3781
      Detection Prevalence : 0.4572
      Balanced Accuracy : 0.7482

      'Positive' Class : 1

Decision Tree Testing Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0  116788  57916
      1   25439 121669

      Accuracy : 0.741
      95% CI : (0.7395, 0.7425)
      No Information Rate : 0.558
      P-Value [Acc > NIR] : < 2.2e-16

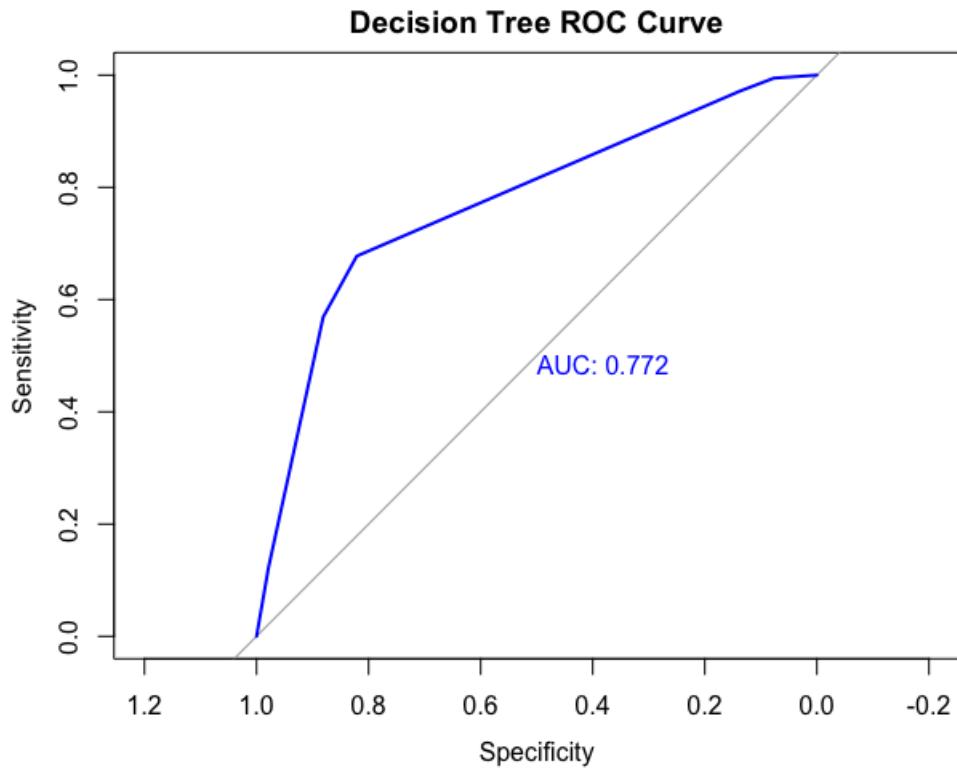
      Kappa : 0.4871

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.6775
      Specificity : 0.8211
      Pos Pred Value : 0.8271
      Neg Pred Value : 0.6685
      Prevalence : 0.5580
      Detection Rate : 0.3781
      Detection Prevalence : 0.4571
      Balanced Accuracy : 0.7493

      'Positive' Class : 1

```



5.5 K Nearest Neighbor

5.5.1 Confusion Matrix

The testing accuracy has decreased by a 13%. As KNN is a low bias high variance model, has overfit on the training set. Altering the k value can solve this problem.

```

KNN Training Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 287876 47639
      1 42907 372471

      Accuracy : 0.8794
      95% CI : (0.8787, 0.8802)
      No Information Rate : 0.5595
      P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.7557

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.8866
      Specificity : 0.8703
      Pos Pred Value : 0.8967
      Neg Pred Value : 0.8580
      Prevalence : 0.5595
      Detection Rate : 0.4960
      Detection Prevalence : 0.5532
      Balanced Accuracy : 0.8784

      'Positive' Class : 1

```

```

KNN Testing Metrics:
Confusion Matrix and Statistics

      Reference
Prediction      0      1
      0 101777 40632
      1 40450 138953

      Accuracy : 0.748
      95% CI : (0.7465, 0.7495)
      No Information Rate : 0.558
      P-Value [Acc > NIR] : <2e-16

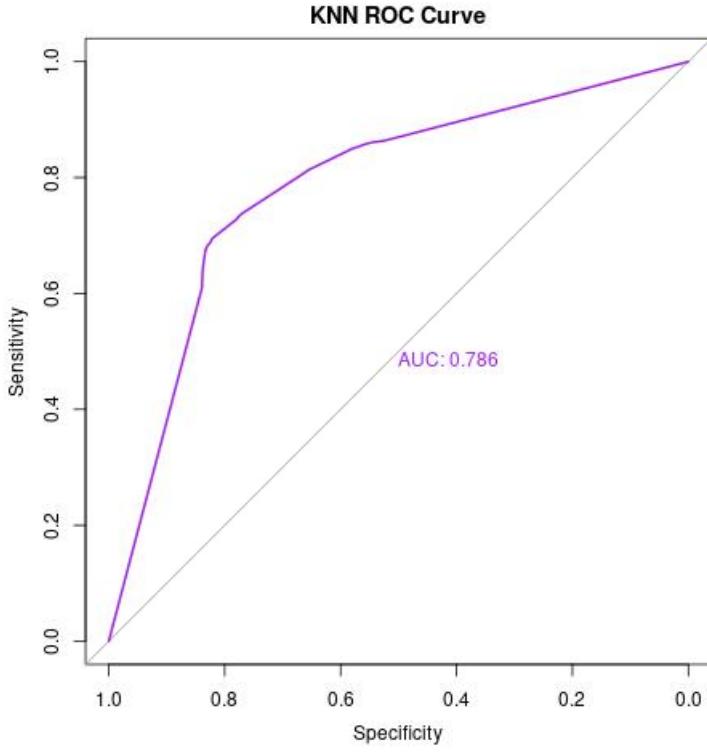
      Kappa : 0.4893

McNemar's Test P-Value : 0.525

      Sensitivity : 0.7737
      Specificity : 0.7156
      Pos Pred Value : 0.7745
      Neg Pred Value : 0.7147
      Prevalence : 0.5580
      Detection Rate : 0.4318
      Detection Prevalence : 0.5575
      Balanced Accuracy : 0.7447

      'Positive' Class : 1

```



5.6 Tabulating Model Performance on Test Set

This table is drawn with respect to model's performance on the testing set.

Model	Accuracy Score	Recall(Sensitivity)	Specificity	Precision	AUC
GLM	0.70	0.66	0.74	0.76	0.77
Naive Bayes	0.61	0.37	0.92	0.85	0.68
SVM	0.74	0.66	0.85	0.85	0.75
Decision Tree	0.74	0.67	0.82	0.82	0.77
KNN	0.74	0.77	0.71	0.77	0.78

5.7 Observations:

- Naive Bayes is the weakest performing model out of all the models. This might be because naive bayes is a high bias classifier. It assumes that all the features are independent of each other. This can lead to underfitting. Naive Bayes model has very low recall score (accuracy to detect +ve samples), as per the application of the dataset we can reject naive bayes model.

- SVM and Decision Tree are the top performing models. SVM and Decision Tree models are high-variance models, meaning large changes to training dataset/features which have high importance may cause the model to under perform as it might not understand underlying mapping between the features as good as say logistic regression.

Chapter 6

Ensemble Techniques

In this chapter we focus on computational part of modeling. Training models on large datasets consumes a lot of time. Techniques like cross validation and stacking can be implemented in parallel in order to reduce modeling time. R libraries used for parallel implementation: `foreach` and `doParallel`. We register all available cores using: `registerDoParallel(detectCores())`.

6.1 K Fold Cross Validation

We implement GLM with K-fold cross validation, where $k = 10$. We compare the model's performance and system performance with parallel and sequential implementation. We also check if GLM performs better with cross validation for our dataset.

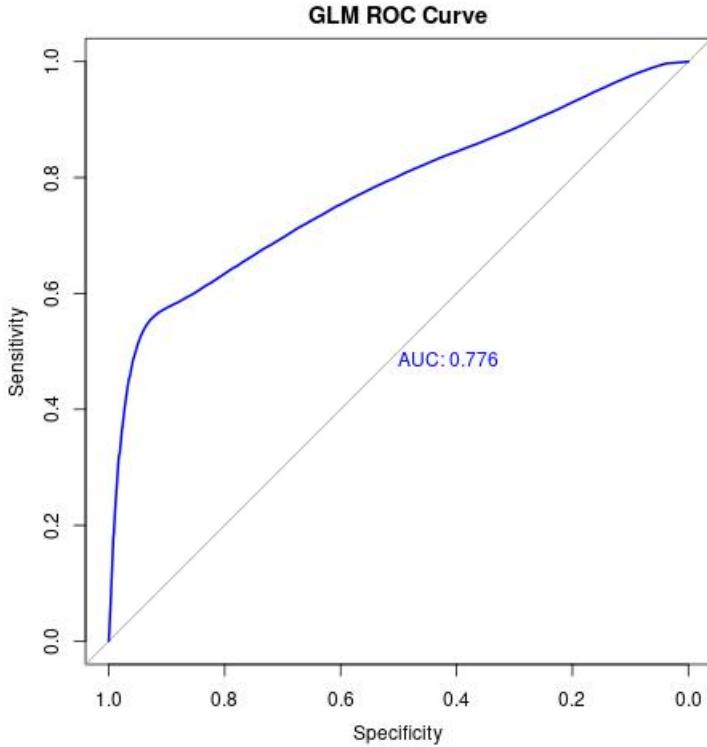
We create 10 folds using: `cv = createFolds(data$Delay, k=10, list=T)`.

GLM with cross validation (parallel).

```
Acc:  
[1] 0.7  
Setting levels: control = 0, case = 1  
Setting direction: controls < cases  
Area under the curve: 0.7761  
    user   system elapsed  
 77.290   4.325 10.071
```

GLM with cross validation (sequential).

```
Acc:  
[1] 0.7  
    user   system elapsed  
 61.847   0.591  62.447
```



While the accuracy remains the same, the parallel implementation takes 8 seconds (elapsed time) while the sequential implementation takes over a minute. This is an incredible advantage when working with large datasets. Performance of GLM is the same with and without cross validation.

6.2 Boosting

Boosting works by sequentially adding predictors to an ensemble, each new predictor corrects the predecessor. Gradient boosting has the highest test accuracy and AUC out of all the classifiers. Being a sequential algorithm, boosting cannot be implemented in parallel, we implement k-fold cross validation with boosting in parallel, which reduces training time while maintaining the performance. GBM with cross validation (parallel 1500 trees).

```

Acc:
[1] 0.76
Setting levels: control = 0, case = 1
Setting direction: controls < cases
Area under the curve: 0.8479
    user     system    elapsed
5817.049     4.419   583.613

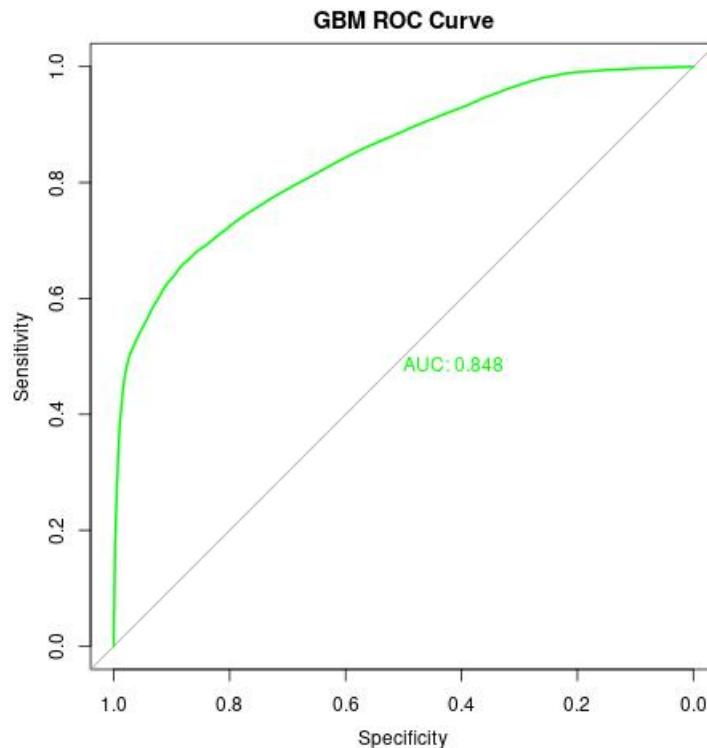
```

GBM with cross validation (sequential 1500 trees).

```

Acc:
[1] 0.76
    user     system    elapsed
4695.444     0.000 4698.581

```



6.3 Random Forest

Random forest applies bagging for training. The training dataset is sampled with replacement and uses majority vote to classify a sample. Bagging helps in reduction of variance without increase in the bias.

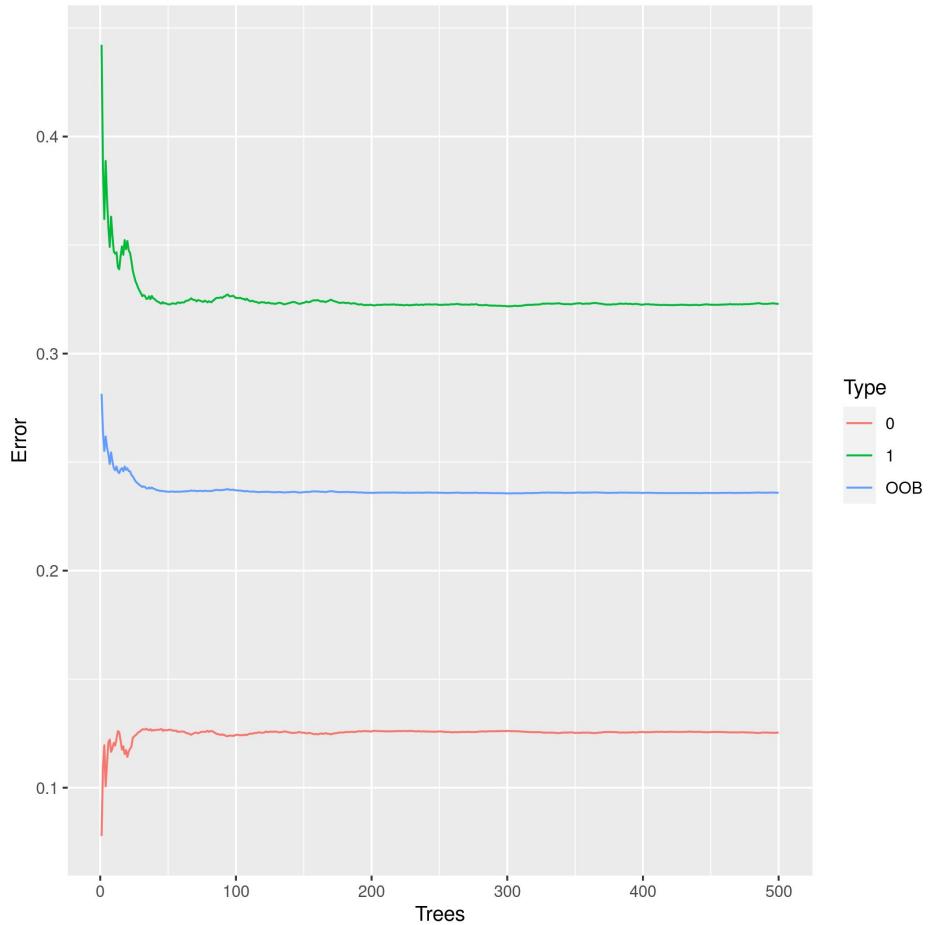
```

Type of random forest: classification
Number of trees: 500
No. of variables tried at each split: 3

OOB estimate of error rate: 23.54%
Confusion matrix:
      0     1 class.error
0 289213 41570  0.1256715
1 135204 284906  0.3218300

```

Out of bag error is calculated on the training set to find optimum number of trees. Out of bag error is the mean error on samples not present in the bootstrapped sample. We observe that the error levels off at round 100 trees for our dataset.



```

Confusion Matrix and Statistics

           Reference
Prediction      0      1
      0 124347 17880
      1  57389 122196

          Accuracy : 0.7661
          95% CI  : (0.7646, 0.7676)
          No Information Rate : 0.5647
          P-Value [Acc > NIR] : < 2.2e-16

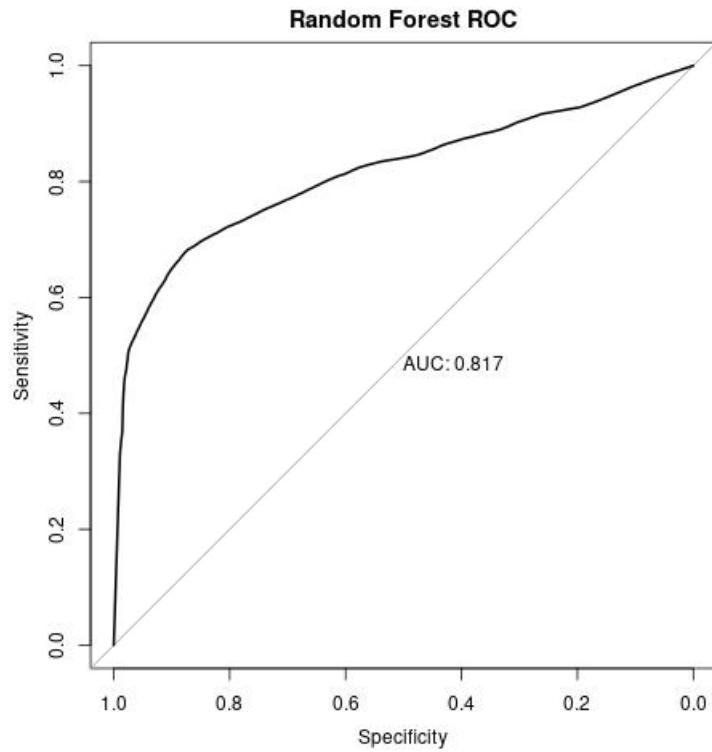
          Kappa : 0.5391

McNemar's Test P-Value : < 2.2e-16

          Sensitivity : 0.8724
          Specificity  : 0.6842
          Pos Pred Value : 0.6804
          Neg Pred Value : 0.8743
          Prevalence   : 0.4353
          Detection Rate : 0.3797
          Detection Prevalence : 0.5580
          Balanced Accuracy : 0.7783

          'Positive' Class : 1

```



Random forest has test accuracy of 76.6% and AUC of 0.81 which is second

highest after gradient boosting.

6.4 Tabulating Ensemble Method Performance

Method	Sequential Time	Parallel Time	Accuracy	AUC
K-Fold Cross Validation	62.45	10.07	0.70	0.77
Gradient Boosting	4698.58	583.61	0.76	0.85
Random Forest	-	-	0.76	0.82

6.5 Conclusion

- Primary motivation to implement ensemble methods on a large dataset is to implement these methods in parallel to reduce the time take to train without giving up on model performance.
- Parallel implementation of gradient boosting model with k-fold cross validation executes 10x faster than sequential implementation while giving the best performance on the test set.
- GBM and random forest are the best suited models for this dataset. Both have high scores metrics on the test set and low computation time.